

Natural language processing has its roots in the 1950s. Already in 1950, [Alan Turing](#) published an article titled "[Computing Machinery and Intelligence](#)" which proposed what is now called the [Turing test](#) as a criterion of intelligence, though at the time that was not articulated as a problem separate from artificial intelligence. The proposed test includes a task that involves the automated interpretation and generation of natural language.

## Symbolic NLP (1950s – early 1990s)[\[edit\]](#)

The premise of symbolic NLP is well-summarized by [John Searle](#)'s [Chinese room](#) experiment: Given a collection of rules (e.g., a Chinese phrasebook, with questions and matching answers), the computer emulates natural language understanding (or other NLP tasks) by applying those rules to the data it confronts.

- **1950s:** The [Georgetown experiment](#) in 1954 involved fully [automatic translation](#) of more than sixty Russian sentences into English. The authors claimed that within three or five years, machine translation would be a solved problem.<sup>[2]</sup> However, real progress was much slower, and after the [ALPAC report](#) in 1966, which found that ten-year-long research had failed to fulfill the expectations, funding for machine translation was dramatically reduced. Little further research in machine translation was conducted until the late 1980s when the first [statistical machine translation](#) systems were developed.
- **1960s:** Some notably successful natural language processing systems developed in the 1960s were [SHRDLU](#), a natural language system working in restricted "[blocks worlds](#)" with restricted vocabularies, and [ELIZA](#), a simulation of a [Rogerian psychotherapist](#), written by [Joseph Weizenbaum](#) between 1964 and 1966. Using almost no information about human thought or emotion, ELIZA sometimes provided a startlingly human-like interaction. When the "patient" exceeded the very small knowledge base, ELIZA might provide a generic response, for example, responding to "My head hurts" with "Why do you say your head hurts?".
- **1970s:** During the 1970s, many programmers began to write "conceptual [ontologies](#)", which structured real-world information into computer-understandable data. Examples are MARGIE (Schank, 1975), SAM (Cullingford, 1978), PAM (Wilensky, 1978), TaleSpin (Meehan, 1976), QUALM (Lehnert, 1977), Politics (Carbonell, 1979), and Plot Units (Lehnert 1981). During this time, the first [chatterbots](#) were written (e.g., [PARRY](#)).
- **1980s:** The 1980s and early 1990s mark the hey-day of symbolic methods in NLP. Focus areas of the time included research on rule-based parsing (e.g., the development of [HPSG](#) as a computational operationalization of [generative grammar](#)), morphology (e.g., two-level morphology<sup>[3]</sup>), semantics (e.g., [Lesk algorithm](#)), reference (e.g., within Centering Theory<sup>[4]</sup>) and other areas of natural language understanding (e.g., in the [Rhetorical Structure Theory](#)). Other lines of research were continued, e.g., the development of chatterbots with [Racter](#) and [Jabberwacky](#). An important development (that eventually led to the statistical turn in the 1990s) was the rising importance of quantitative evaluation in this period.<sup>[5]</sup>

## Statistical NLP (1990s–2010s)[\[edit\]](#)

Up to the 1980s, most natural language processing systems were based on complex sets of hand-written rules. Starting in the late 1980s, however, there was a revolution in natural language processing with the introduction of [machine learning](#) algorithms for language processing. This was due to both the steady increase in computational power (see [Moore's law](#)) and the gradual lessening of the dominance of [Chomskyan](#) theories of linguistics (e.g. [transformational grammar](#)), whose theoretical underpinnings discouraged the sort of [corpus linguistics](#) that underlies the machine-learning approach to language processing.<sup>[6]</sup>

- **1990s:** Many of the notable early successes on statistical methods in NLP occurred in the field of [machine translation](#), due especially to work at IBM Research. These systems were able to take advantage of existing multilingual [textual corpora](#) that had

been produced by the [Parliament of Canada](#) and the [European Union](#) as a result of laws calling for the translation of all governmental proceedings into all official languages of the corresponding systems of government. However, most other systems depended on corpora specifically developed for the tasks implemented by these systems, which was (and often continues to be) a major limitation in the success of these systems. As a result, a great deal of research has gone into methods of more effectively learning from limited amounts of data.

- **2000s:** With the growth of the web, increasing amounts of raw (unannotated) language data has become available since the mid-1990s. Research has thus increasingly focused on [unsupervised](#) and [semi-supervised learning](#) algorithms. Such algorithms can learn from data that has not been hand-annotated with the desired answers or using a combination of annotated and non-annotated data. Generally, this task is much more difficult than [supervised learning](#), and typically produces less accurate results for a given amount of input data. However, there is an enormous amount of non-annotated data available (including, among other things, the entire content of the [World Wide Web](#)), which can often make up for the inferior results if the algorithm used has a low enough [time complexity](#) to be practical.

## Neural NLP (present)[\[edit\]](#)

In the 2010s, [representation learning](#) and [deep neural network](#)-style machine learning methods became widespread in natural language processing. That popularity was due partly to a flurry of results showing that such techniques<sup>[7][8]</sup> can achieve state-of-the-art results in many natural language tasks, e.g., in [language modeling](#)<sup>[9]</sup> and parsing.<sup>[10][11]</sup> This is increasingly important in [medicine](#) and [healthcare](#), where NLP helps analyze notes and text in [electronic health records](#) that would otherwise be inaccessible for study when seeking to improve care.<sup>[12]</sup>

## Methods: Rules, statistics, neural networks[\[edit\]](#)

---

In the early days, many language-processing systems were designed by symbolic methods, i.e., the hand-coding of a set of rules, coupled with a dictionary lookup:<sup>[13][14]</sup> such as by writing grammars or devising heuristic rules for [stemming](#).

More recent systems based on [machine-learning](#) algorithms have many advantages over hand-produced rules:

- The learning procedures used during machine learning automatically focus on the most common cases, whereas when writing rules by hand it is often not at all obvious where the effort should be directed.
- Automatic learning procedures can make use of statistical inference algorithms to produce models that are robust to unfamiliar input (e.g. containing words or structures that have not been seen before) and to erroneous input (e.g. with misspelled words or words accidentally omitted). Generally, handling such input gracefully with handwritten rules, or, more generally, creating systems of handwritten rules that make soft decisions, is extremely difficult, error-prone and time-consuming.
- Systems based on automatically learning the rules can be made more accurate simply by supplying more input data. However, systems based on handwritten rules can only be made more accurate by increasing the complexity of the rules, which is a much more difficult task. In particular, there is a limit to the complexity of systems based on handwritten rules, beyond which the systems become more and more unmanageable. However, creating more data to input to machine-learning systems simply requires a corresponding increase in the number of man-hours worked, generally without significant increases in the complexity of the annotation process.

Despite the popularity of machine learning in NLP research, symbolic methods are still (2020) commonly used:

- when the amount of training data is insufficient to successfully apply machine learning methods, e.g., for the machine translation of low-resource languages such as provided by the [Apertium](#) system,
- for preprocessing in NLP pipelines, e.g., [tokenization](#), or
- for postprocessing and transforming the output of NLP pipelines, e.g., for [knowledge extraction](#) from syntactic parses.

## **Statistical methods**